

CS180 Notes 3-3-15

- Thread Libraries
 - Similar to processes having parent/child, threads have main/worker relationships
 - Threads are managed with three methods: *pthread_create*, *pthread_join*, and *pthread_exit*.
 - *pthread_create* has four parameters
 - A *pthread_t* ID
 - Specifiers for the thread
 - The new function to be called
 - Parameters to the new function
 - If multiple parameters, must pass as a struct
 - *pthread_exit* can be an expensive call, so sometimes it is best to let the OS reclaim the thread
 - Valgrind can provide a false memory leak with *pthread_exit*.
 - As always, check the return values!
- All thread functions have the same type of signature
 - *void *[name](void *p)*
 - Assists with being generic
 - Does not return to the caller
- All global data is shared between all threads
 - No extra work to share data
 - Safe for reads, but not safe for writing.
- Thread IDs tend to be very long, so that's why we are giving the tid by 1000000 to have more reasonable IDs to print.
- While we are waiting with *pthread_join*, the thread is now blocked until that thread completes.
 - It doesn't matter if other threads finish first. We still wait for the one we designate.
 - After a join returns, the thread is terminated.
- We have the same process ID as threads are inside processes.
 - Thread IDs are (effectively) unique.
 - There is absolutely no way to control what thread runs at a given time.
 - There is a way to make a thread idle.
- In production code, random number generators are not considered thread safe.
- Example where threads are insanely useful: File search on a hard disk
 - Create threads to search each drive in the computer
 - No point in putting multiple threads per drive – the I/O is already the bottleneck
- A way to avoid using global data is to communicate using data structs.

- When using a thread's parameter, we must cast the data into the proper type we want.
 - Example: `(data_struct *)data` (from `void *data`)
- Since we're scoped into main, we must pass array data by structs (see example in notes)
 - Because the data struct is the name for all threads, we could have passed the one struct to each thread as opposed to a struct for each thread.
 - It is potentially worth it to check for type of data passed into the thread's new function for safety.
- Each thread is passed a pointer to its own data
 - Threads need to know what data to expect.
 - The data must still remain alive while other threads are potentially using it
- **stdio** library functions are thread safe with GNU and Microsoft compilers
 - Related to the concept of re-entry. (Executing the same code at the same time)
 - *localtime* is not thread safe, but *localtime_t* is.
 - *strace* is not thread safe as static structures are used
 - We can still get interleaved output, but nothing will be corrupted.
- Circle, Leibniz, and atan methods
 - Before, we had these functions running serially – one after another.
 - Call, wait to finish, call the next one, wait, etc.
 - With threads (and processes), we can all three functions concurrently, which drastically reduces the run-time of our program.
 - The user time is still the cumulative time of all the process put together.
 - Remember that threads share global data. Fortunately, each method only writes to its own global data, so we are okay.